

# Hidden Markov Models for Regime Detection

Technical Implementation Guide

Precog Quant Task 2026

Technical Documentation

February 9, 2026

## Contents

# 1 Introduction

Hidden Markov Models (HMMs) provide a principled framework for identifying latent market regimes from observable price dynamics. Rather than predicting returns directly, HMMs excel at:

- **Regime Identification:** Distinguishing bull/bear, high/low volatility states
- **Adaptive Conditioning:** Modifying signals based on regime
- **Risk Management:** Adjusting position sizing by regime
- **Feature Generation:** Creating regime-aware ML features

## 1.1 Why HMM for Trading?

Financial markets exhibit regime-switching behavior:

- Trending vs. mean-reverting periods
- High volatility vs. low volatility regimes
- Risk-on vs. risk-off environments

HMMs capture this through:

- Discrete latent states representing regimes
- State-dependent emission distributions
- Markovian transition dynamics


# 2 Mathematical Foundation

## 2.1 HMM Definition

**Definition 1** (Hidden Markov Model). A Hidden Markov Model with  $K$  states consists of:

- **State space:**  $\mathcal{S} = \{1, 2, \dots, K\}$
- **Initial distribution:**  $\pi = (\pi_1, \dots, \pi_K)$  where  $\pi_k = P(S_1 = k)$
- **Transition matrix:**  $\mathbf{A} = [a_{jk}]$  where  $a_{jk} = P(S_t = k | S_{t-1} = j)$
- **Emission distributions:**  $B = \{b_k(y)\}$  where  $b_k(y) = P(Y_t = y | S_t = k)$

## 2.2 Graphical Model



## 2.3 Gaussian HMM

For financial applications, we use Gaussian emissions:

$$b_k(y) = \frac{1}{\sqrt{2\pi\sigma_k^2}} \exp\left(-\frac{(y - \mu_k)^2}{2\sigma_k^2}\right) \quad (1)$$

The parameters are:

- $\mu_k$ : Mean return/feature in state  $k$
- $\sigma_k$ : Volatility in state  $k$

**Remark 1** (State Interpretation). For a 2-state volatility HMM:

- State 1 (Low Vol):  $\sigma_1 < \sigma_2$  (calm markets)
- State 2 (High Vol):  $\sigma_2 > \sigma_1$  (turbulent markets)

## 3 HMM Algorithms

### 3.1 The Three Fundamental Problems

1. **Evaluation:** Given model  $\lambda = (\mathbf{A}, \mathbf{B}, \boldsymbol{\pi})$  and observations  $\mathbf{Y} = (y_1, \dots, y_T)$ , compute  $P(\mathbf{Y}|\lambda)$
2. **Decoding:** Find the most likely state sequence  $\mathbf{S}^* = \arg \max_{\mathbf{S}} P(\mathbf{S}|\mathbf{Y}, \lambda)$
3. **Learning:** Estimate parameters  $\hat{\lambda} = \arg \max_{\lambda} P(\mathbf{Y}|\lambda)$

### 3.2 Forward Algorithm

Computes  $\alpha_t(k) = P(y_1, \dots, y_t, S_t = k)$  recursively:

---

#### Algorithm 1 Forward Algorithm

---

- 1: **Initialize:**  $\alpha_1(k) = \pi_k \cdot b_k(y_1)$  for all  $k$
  - 2: **for**  $t = 2$  to  $T$  **do**
  - 3:   **for**  $k = 1$  to  $K$  **do**
  - 4:      $\alpha_t(k) = b_k(y_t) \sum_{j=1}^K \alpha_{t-1}(j) \cdot a_{jk}$
  - 5:   **end for**
  - 6: **end for**
  - 7: **Likelihood:**  $P(\mathbf{Y}|\lambda) = \sum_{k=1}^K \alpha_T(k)$
- 

### 3.3 Backward Algorithm

Computes  $\beta_t(k) = P(y_{t+1}, \dots, y_T | S_t = k)$ :

---

#### Algorithm 2 Backward Algorithm

---

- 1: **Initialize:**  $\beta_T(k) = 1$  for all  $k$
  - 2: **for**  $t = T - 1$  down to 1 **do**
  - 3:   **for**  $j = 1$  to  $K$  **do**
  - 4:      $\beta_t(j) = \sum_{k=1}^K a_{jk} \cdot b_k(y_{t+1}) \cdot \beta_{t+1}(k)$
  - 5:   **end for**
  - 6: **end for**
-

### 3.4 State Probabilities

**Filtering** (real-time, no future info):

$$P(S_t = k | y_1, \dots, y_t) = \frac{\alpha_t(k)}{\sum_{j=1}^K \alpha_t(j)} \quad (2)$$

**Smoothing** (uses all data, look-ahead!):

$$\gamma_t(k) = P(S_t = k | \mathbf{Y}) = \frac{\alpha_t(k) \cdot \beta_t(k)}{\sum_{j=1}^K \alpha_t(j) \cdot \beta_t(j)} \quad (3)$$

**Remark 2** (Critical Warning). **Filtering** uses only past data  $\Rightarrow$  Valid for OOS  
**Smoothing** uses future data  $\Rightarrow$  IS only (look-ahead bias!)

### 3.5 Viterbi Algorithm

Finds the most likely state sequence (decoding):

---

#### Algorithm 3 Viterbi Algorithm

---

- 1: **Initialize:**  $\delta_1(k) = \pi_k \cdot b_k(y_1)$ ,  $\psi_1(k) = 0$
  - 2: **for**  $t = 2$  to  $T$  **do**
  - 3:     **for**  $k = 1$  to  $K$  **do**
  - 4:          $\delta_t(k) = \max_j [\delta_{t-1}(j) \cdot a_{jk}] \cdot b_k(y_t)$
  - 5:          $\psi_t(k) = \arg \max_j [\delta_{t-1}(j) \cdot a_{jk}]$
  - 6:     **end for**
  - 7: **end for**
  - 8: **Termination:**  $S_T^* = \arg \max_k \delta_T(k)$
  - 9: **Backtrack:**  $S_t^* = \psi_{t+1}(S_{t+1}^*)$  for  $t = T - 1, \dots, 1$
- 

## 4 Parameter Estimation: Baum-Welch (EM)

### 4.1 Overview

The Baum-Welch algorithm is EM specialized for HMMs.

### 4.2 E-Step: Compute Expected Statistics

Define transition probability:

$$\xi_t(j, k) = P(S_t = j, S_{t+1} = k | \mathbf{Y}) = \frac{\alpha_t(j) \cdot a_{jk} \cdot b_k(y_{t+1}) \cdot \beta_{t+1}(k)}{P(\mathbf{Y})} \quad (4)$$

### 4.3 M-Step: Update Parameters

$$\hat{\pi}_k = \gamma_1(k) \quad (5)$$

$$\hat{a}_{jk} = \frac{\sum_{t=1}^{T-1} \xi_t(j, k)}{\sum_{t=1}^{T-1} \gamma_t(j)} \quad (6)$$

$$\hat{\mu}_k = \frac{\sum_{t=1}^T \gamma_t(k) \cdot y_t}{\sum_{t=1}^T \gamma_t(k)} \quad (7)$$

$$\hat{\sigma}_k^2 = \frac{\sum_{t=1}^T \gamma_t(k) \cdot (y_t - \hat{\mu}_k)^2}{\sum_{t=1}^T \gamma_t(k)} \quad (8)$$

## 5 Implementation

### 5.1 Python Implementation

```

1 import numpy as np
2 from scipy.stats import norm
3
4 class GaussianHMM:
5     """Gaussian Hidden Markov Model for regime detection."""
6
7     def __init__(self, n_states=2, n_iter=100, tol=1e-6):
8         self.n_states = n_states
9         self.n_iter = n_iter
10        self.tol = tol
11
12    def _initialize(self, y):
13        """Initialize parameters via k-means-like clustering."""
14        T = len(y)
15        K = self.n_states
16
17        # Initialize by quantiles
18        quantiles = np.percentile(y, np.linspace(0, 100, K+1))
19        self.means_ = np.array([
20            np.mean(y[(y >= quantiles[k]) & (y < quantiles[k+1])])
21            for k in range(K)
22        ])
23        self.vars_ = np.var(y) * np.ones(K)
24
25        # Transition matrix (sticky)
26        self.transmat_ = np.full((K, K), 0.1 / (K-1))
27        np.fill_diagonal(self.transmat_, 0.9)
28
29        # Initial distribution
30        self.startprob_ = np.ones(K) / K
31
32    def _forward(self, y):
33        """Forward algorithm (log-space for numerical stability)."""
34        T = len(y)
35        K = self.n_states
36
37        log_alpha = np.zeros((T, K))
38
39        # Initialize
40        for k in range(K):
41            log_alpha[0, k] = np.log(self.startprob_[k]) + \
42                norm.logpdf(y[0], self.means_[k], np.sqrt(
43                    self.vars_[k]))
44
45        # Recurse
46        for t in range(1, T):
47            for k in range(K):
48                log_sum = np.logaddexp.reduce(
49                    log_alpha[t-1, :] + np.log(self.transmat_[ :, k])
50                )
51                log_alpha[t, k] = log_sum + \
52                    norm.logpdf(y[t], self.means_[k], np.

```

```

53     return log_alpha
54
55     def _backward(self, y):
56         """Backward algorithm (log-space)."""
57         T = len(y)
58         K = self.n_states
59
60         log_beta = np.zeros((T, K))
61         # log_beta[T-1, :] = 0 (log(1) = 0)
62
63         for t in range(T-2, -1, -1):
64             for j in range(K):
65                 terms = []
66                 for k in range(K):
67                     term = np.log(self.transmat_[j, k]) + \
68                         norm.logpdf(y[t+1], self.means_[k], np.sqrt(
self.vars_[k])) + \
69                         log_beta[t+1, k]
70                     terms.append(term)
71                 log_beta[t, j] = np.logaddexp.reduce(terms)
72
73         return log_beta
74
75     def fit(self, y):
76         """Fit HMM via Baum-Welch (EM)."""
77         self._initialize(y)
78         T = len(y)
79         K = self.n_states
80
81         prev_log_lik = -np.inf
82
83         for iteration in range(self.n_iter):
84             # E-step
85             log_alpha = self._forward(y)
86             log_beta = self._backward(y)
87
88             # Log-likelihood
89             log_lik = np.logaddexp.reduce(log_alpha[-1, :])
90
91             # Gamma (state probabilities)
92             log_gamma = log_alpha + log_beta
93             log_gamma -= np.logaddexp.reduce(log_gamma, axis=1,
keepdims=True)
94             gamma = np.exp(log_gamma)
95
96             # Xi (transition probabilities)
97             log_xi = np.zeros((T-1, K, K))
98             for t in range(T-1):
99                 for j in range(K):
100                     for k in range(K):
101                         log_xi[t, j, k] = log_alpha[t, j] + \
102 + \
103                         norm.logpdf(y[t+1], self.
means_[k], np.sqrt(self.vars_[k])) + \
104                         log_beta[t+1, k]
105                 log_xi[t] -= np.logaddexp.reduce(log_xi[t].flatten())
106             xi = np.exp(log_xi)

```

```

107
108     # M-step
109     self.startprob_ = gamma[0] + 1e-10
110     self.startprob_ /= self.startprob_.sum()
111
112     for j in range(K):
113         for k in range(K):
114             self.transmat_[j, k] = xi[:, j, k].sum() / gamma
115             self.transmat_ += 1e-10
116             self.transmat_ /= self.transmat_.sum(axis=1, keepdims=True)
117
118         for k in range(K):
119             gamma_k = gamma[:, k]
120             self.means_[k] = np.sum(gamma_k * y) / np.sum(gamma_k)
121             self.vars_[k] = np.sum(gamma_k * (y - self.means_[k])
122             **2) / np.sum(gamma_k)
123             self.vars_[k] = max(self.vars_[k], 1e-6) # Prevent
124             zero variance
125
126     # Convergence
127     if abs(log_lik - prev_log_lik) < self.tol:
128         break
129     prev_log_lik = log_lik
130
131     # Reorder states by volatility
132     order = np.argsort(self.vars_)
133     self.means_ = self.means_[order]
134     self.vars_ = self.vars_[order]
135     self.transmat_ = self.transmat_[order][:, order]
136     self.startprob_ = self.startprob_[order]
137
138     return self
139
140 def predict_proba(self, y):
141     """Predict state probabilities (filtering - no future info!).
142     """
143     log_alpha = self._forward(y)
144     log_proba = log_alpha - np.logaddexp.reduce(log_alpha, axis=1,
145     keepdims=True)
146     return np.exp(log_proba)
147
148 def predict(self, y):
149     """Predict most likely states (Viterbi)."""
150     T = len(y)
151     K = self.n_states
152
153     log_delta = np.zeros((T, K))
154     psi = np.zeros((T, K), dtype=int)
155
156     # Initialize
157     for k in range(K):
158         log_delta[0, k] = np.log(self.startprob_[k]) + \
159         norm.logpdf(y[0], self.means_[k], np.sqrt(
160         self.vars_[k]))
161
162     # Recurse
163     for t in range(1, T):

```

```

159         for k in range(K):
160             candidates = log_delta[t-1, :] + np.log(self.transmat_
161            [:, k])
162             psi[t, k] = np.argmax(candidates)
163             log_delta[t, k] = candidates[psi[t, k]] + \
164                 norm.logpdf(y[t], self.means_[k], np.
165                 sqrt(self.vars_[k]))
166
167         # Backtrack
168         states = np.zeros(T, dtype=int)
169         states[-1] = np.argmax(log_delta[-1, :])
170         for t in range(T-2, -1, -1):
171             states[t] = psi[t+1, states[t+1]]
172
173         return states

```

Listing 1: Gaussian HMM Class

## 6 Feature Generation for ML Models

### 6.1 HMM-Derived Features

Feature	Formula	Interpretation
hmm_state	$\arg \max_k P(S_t = k   y_{1:t})$	Most likely regime
hmm_prob_high	$P(S_t = \text{high\_vol}   y_{1:t})$	High-vol regime probability
hmm_entropy	$-\sum_k p_k \log p_k$	Regime uncertainty
hmm_transition	$1[S_t \neq S_{t-1}]$	Regime change indicator
hmm_stability	Rolling mean of hmm_transition	Regime persistence

Table 1: HMM-Derived ML Features

### 6.2 Feature Usage Strategy

```

1 def generate_hmm_features(hmm, observations):
2     """Generate ML features from fitted HMM."""
3     # State probabilities (filtering - valid for OOS)
4     proba = hmm.predict_proba(observations)
5
6     # Most likely state
7     states = hmm.predict(observations)
8
9     features = {
10         'hmm_state': states,
11         'hmm_prob_high_vol': proba[:, -1], # Last state = highest vol
12         'hmm_prob_low_vol': proba[:, 0],   # First state = lowest vol
13         'hmm_entropy': -np.sum(proba * np.log(proba + 1e-10), axis=1),
14         'hmm_transition': np.concatenate([[0], np.diff(states) != 0]).
15         astype(float),
16     }
17
18     # Regime stability (rolling 21-day)
19     features['hmm_stability'] = pd.Series(features['hmm_transition']).
20     rolling(21).mean().values

```



```
20 return pd.DataFrame(features)
```

Listing 2: HMM Feature Generation

## 7 Critical Implementation Rules

### 7.1 In-Sample vs Out-of-Sample Usage

Component	IS	OOS	Reasoning
Baum-Welch Training	✓	×	Freeze params from IS
Forward (Filtering)	✓	✓	Only uses past data
Backward (Smoothing)	✓	×	Uses future data
Viterbi (Global)	✓	×	Uses all observations
Real-time Filtering	✓	✓	Valid for production

Table 2: HMM Component Usage Rules

### 7.2 State Label Consistency

**Remark 3** (Label Flipping Problem). HMM state labels are arbitrary — training on different subsets may swap state meanings!

**Solution:** Order states by a consistent criterion:

- By variance: State 0 = lowest vol, State K-1 = highest vol
- By mean: State 0 = lowest return, etc.

### 7.3 Validation Checks

1. **State Interpretability:** Verify states have distinct  $\mu_k$  or  $\sigma_k$
2. **Transition Stability:** Check  $a_{jj} \gg a_{jk}$  (sticky states)
3. **IS/OOS Consistency:** Compare state distributions across periods

## 8 Application in Our Pipeline

### 8.1 Philosophy

*“Use HMM as a regime filter, NOT a return predictor.”*

HMMs improve signal stability and risk management, not alpha magnitude.

### 8.2 Training Phase (In-Sample: 2016-2023)

1. Select input features (e.g., volatility, filtered returns)
2. Fit HMM with 2-4 states using Baum-Welch
3. Order states by volatility for consistent labeling
4. Generate regime features via filtering

### 8.3 Production Phase (Out-of-Sample: 2024-2026)

1. Use frozen HMM parameters from IS
2. Apply forward algorithm for real-time filtering
3. Generate regime features
4. Condition trading signals on regime state

### 8.4 Signal Conditioning Example

```

1 def regime_adjusted_signal(raw_signal, hmm_proba, scale_factors=[1.0,
2   0.5]):
3     """Scale signals based on regime."""
4     # scale_factors: [low_vol_scale, high_vol_scale]
5     weights = hmm_proba @ np.array(scale_factors)
6     return raw_signal * weights

```

Listing 3: Regime-Conditional Trading

## 9 Empirical Results

### 9.1 Regime Characteristics (2-State Model)

Statistic	Low Vol Regime	High Vol Regime
Mean Return (%/day)	~0.05	~-0.02
Volatility (%/day)	~0.8	~2.5
Duration (days avg.)	~45	~15
Frequency (%)	~75	~25

Table 3: Typical 2-State HMM Regime Statistics

### 9.2 Impact on Strategy

- **Without HMM:** IS Sharpe 2.14, OOS Sharpe 2.10
- **With HMM Features:** IS Sharpe 2.14, OOS Sharpe 2.19

HMM improved stability but not raw performance — exactly as expected.

## 10 Common Pitfalls

1. **Using Smoothed Probabilities OOS**
  - Smoothing ( $\gamma_t$ ) uses future data
  - Use filtering ( $\alpha_t$  normalized) for OOS
2. **Refitting HMM on OOS Data**
  - Parameters should be frozen from IS training
  - Refitting introduces forward-looking bias
3. **State Label Drift**

- States may flip meaning between runs
- Always reorder by consistent criterion (e.g., variance)

#### 4. Overfitting Number of States

- More states = better IS fit, worse OOS
- Use BIC/AIC for model selection

#### 5. Expecting Direct Alpha

- HMM predicts regimes, not returns
- Use as filter/conditioner, not signal

## 11 Conclusion

Hidden Markov Models provide:

- **Regime identification** from latent market states
- **Principled uncertainty** via state probabilities
- **Feature generation** for ML models
- **Signal conditioning** for risk management

Key implementation rules:

- Train on IS data only, freeze parameters for OOS
- Use forward algorithm (filtering) for production
- Order states consistently (e.g., by volatility)
- Treat HMM as filter, not predictor

*“HMM features improve signal stability, not signal strength — and that’s exactly their value.”*